

TSCC Bootcamp: Introduction to Accessing and Running Jobs on the TSCC System

Compiling & Linking

By: Mary Thomas

TSCC Compilers

Package	Topic Area	License Type	Package Home Page	User Install Location	Installed on: (L)ogin, (C)ompute, (B)oth
CMake	Cross Platform Makefile Generator	Open	CMake Home Page	/opt/cmake	B
GNU Compilers	C and Fortran Compilers	CentOS Core License	GNU Compiler Collection Home Page	/usr/bin/gcc —and— /usr/bin/gfortran	B
Intel Compilers	C and Fortran Compilers	Licensed (flexlm)	Intel Compilers Home Page	/opt/intel	B
Java	Compiler	Open	Java Home Page	/usr/bin/javac	B
PGI Compilers	C and Fortran Compilers	Licensed (flexlm)	PGI Compilers Home Page	/opt/pgi	B
UPC (Unified Parallel C) Compiler	Parallel Computing	BSD	UPC Home Page	/opt/upc	B

Porting Existing Codes

- Copy your application source files to \$HOME directory or to a Data Oasis area for your account
 - /oasis/tscc/scratch/<username> where <username> is your TSCC login name.
- **Data Oasis storage is not backed up**
 - files stored on this system may be lost or destroyed
 - \$HOME directory or project storage are permanent.
- Never run parallel or large serial jobs from the command line – they should always be run using the queue.

Compilers

- **TSCC compute nodes support several parallel programming models:**
 - **MPI:** Default Intel Compiler and openmpi_iblIntel MPI
 - **OpenMP:**
 - All compilers (GNU, Intel, PGI) have OpenMP flags.
 - **GPU nodes:** support CUDA, OpenACC.
 - **Hybrid modes** are possible (see examples below).
- **See:**
 - https://www.sdsc.edu/support/user_guides/tscc.html#compiling
 - https://www.sdsc.edu/support/user_guides/tscc-quick-start.html#gpu-queue

Compiling MPI Codes

- MPI source code should be compiled for TSCC using these commands:
 - **mpicc** [options] file.c (C and C++) myrinet/mx switch and the **Portland Compiler**
 - **mpif77** [options] file.f (Fortran 77 source code) myrinet/mx switch and the **Portland Compiler**
 - **mpif90** [options] file.f (free format code/dynamic memory allocation/object oriented Fortran source code) openmpi and the **Intel compiler**

Compiling Serial Code

- Serial source code should be recompiled
- Portland Group Compilers
 - pgcc [options] file.c (C and C++)
 - pgf77 [options] file.f (fixed form Fortran source code)
 - pgf90 [options] file.f90 (free format Fortran source code)
- Intel Compilers
 - icc [options] file.c (C and C++)
 - ifort [options] file.f (fixed form Fortran source code)
 - ifort [options] file.f90 (free format Fortran source code)
- Gnu Compilers
 - gcc [options] file.c (C and C++)
 - gfortran [options] file.f90 (free format Fortran source code)

To Access Infiniband

- PGI compilers + openmpi_ib
- PGI compilers + mvapich2_ib
- Intel compilers + openmpi_ib
- Intel compilers + mvapich2_ib
- GNU compilers + openmpi_ib
- GNU compilers + mvapich2_ib

Compatibility Options for Fortran

-Mcpp	run the Fortran preprocessor on source files prior to compilation
-Dname[=value]	specify name as a definition to use with conditional compilation directives or the Fortran preprocessor (-Mcpp)
-byteswapio	swap bytes from big-endian to little-endian or vice-versa for unformatted files
-i8	set size of INTEGER and LOGICAL variables to 8 bytes
-i4	set size of INTEGER and LOGICAL variables to 4 bytes
-i2	set size of INTEGER variables to 2 bytes
-r8	treat REAL and CMPLX types as REAL*8 and DCMPLX
-Msave	save all local variables between calls (static allocation)

Detecting Programming Errors

- -g produce symbolic debug information in object file (implies -O0); required for debugging with DDT
- -C/-Mbounds array bounds checking
- -kTrap=(option, option...) specifies behavior on floating point exceptions (used only for main program):
 - dnorm trap on denormalized (very, very small) operands
 - divz trap on divide by zero
 - fp trap on floating point exceptions
 - inexact trap on inexact result
 - inv trap on invalid operands
 - none (default) disables all traps
 - ovf trap on floating point overflow
 - unf trap on floating point underflow
- -traceback add debug information

Optimization levels of Portland Group compilers

- -O0 No optimization
- -O1 (default) Task scheduling within extended basic blocks.
 - Some register allocation; no global optimizations.
- -O2 all level 1 optimizations and global scalar optimizations
 - optimization over all blocks
- -O3/O4 all level 1 and 2 optimizations + more aggressive optimizations.
- -Mflushz flush very, very small values to zero.

PGI Optimization: use **-fast** flag

- Portland Compiler User Guide recommends the use of the **-fast** flag;
- Option is host dependent and has the following default config:
 - set optimization level at -O2
 - unroll loops (-Munroll=c:1)
 - do not generate code to set up a stack frame pointer for every function (-Mnoframe)
 - enable loop redundancy elimination

Using Numerical Libraries: PGI

- PGI includes Optimized ACML library (LAPACK/BLAS/FFT).
- To link to the libraries:

```
pg90/pgf77 myprog.f -llapack -lblas
```

Using Numerical Libraries: Intel

- Intel developed the Math Kernel Library (MKL)
- Linear algebra, FFT and other numerical routines.
 - Basic linear algebra subprograms (BLAS) + sparse routines
 - Fast Fourier Transforms (FFT) in 1 and 2 dimensions, complex and real
 - The linear algebra package, LAPACK
 - A C interface to BLAS
 - Vector Math Library (VML)
 - Vector Statistical Library (VSL)
 - Multi-dimensional Discrete Fourier Transforms (DFTs)
- Documentation available in HTML and PDF format:
 - \${MKL_ROOT}/../Documentation.

Example: Loading the Intel Compilers

- The Intel compilers and the MVAPICH2 MPI implementation will be loaded by default.
- If you have modified your environment, you can reload by executing the module purge & load commands at the Linux prompt, or placing the load commands in your startup file (~/.cshrc or ~/.bashrc)

```
$ module purge
$ module list
No Modulefiles Currently Loaded.
$ module load gnutools
$ module list
Currently Loaded Modulefiles:
 1) gnutools/xx
$ module load intel mvapich2_ib
$ module list
Currently Loaded Modulefiles:
 1) gnutools/xx          2) intel/xx          3) mvapich2_ib/xx
```

Using the PGI Compiler

- PGI (formerly The Portland Group, Inc.), was a company that produced a set of commercially available Fortran, C and C++ compilers for high-performance computing systems.
- It is now owned by NVIDIA.
- PGI compilers can be loaded by executing the following commands at the Linux prompt or placing in your startup file (~/.cshrc or ~/.bashrc).
- For AVX support, compile with -fast

```
[$] module purge
[$] module load gnutools
[$] module load pgi
[$] module list
Currently Loaded Modulefiles:
 1) gnutools/xx    2) pgi/xx
```

- For more information on the PGI compilers run: **man [pgf90 | pgcc | pgCC]**

Parallel Code: MPI Hello World

Change to the MPI examples directory:

```
! Fortran example
program hello
include 'mpif.h'
integer rank, size, ierror, tag, status(MPI_STATUS_SIZE)

call MPI_INIT(ierror)
call MPI_COMM_SIZE(MPI_COMM_WORLD, size, ierror)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierror)
print*, 'node', rank, ': Hello and Welcome to Webinar
Participants!'
call MPI_FINALIZE(ierror)
end
```

Compile the code using :

```
mpif90 -o hello_mpi hello_mpi.f90
```

OpenMP Hello World

```
[PROGRAM OMPHELLO
    INTEGER TNUMBER
    INTEGER OMP_GET_THREAD_NUM

!$OMP PARALLEL DEFAULT(PRIVATE)
    TNUMBER = OMP_GET_THREAD_NUM()
    PRINT *, 'Hello from Thread Number[' , TNUMBER, '] and Welcome Bootcampers!'
!$OMP END PARALLEL

    STOP
END]
```

Check the environment

```
[$] module list
Currently Loaded Modulefiles:
 1) gnutools/2.69          2) pgi.x
```

Compile the code

```
[$] mpif90 -mp -o hello_openmp hello_openmp.f90
```

Hybrid MPI + OpenMP Hello World

```
#include <stdio.h>
#include "mpi.h"
#include <omp.h>

int main(int argc, char *argv[]) {
    int numprocs, rank, namelen;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int iam = 0, np = 1;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Get_processor_name(processor_name, &namelen);

    #pragma omp parallel default(shared) private(iam, np)
    {
        np = omp_get_num_threads();
        iam = omp_get_thread_num();
        printf("Hello Webinar participants from thread %d out of %d from process %d out of %d on %s\n",
               iam, np, rank, numprocs, processor_name);
    }

    MPI_Finalize();
}
```

Compile the code using :

```
Mpif90 -mp -o hello_mpi hello_mpi.f90
```

GPU COMPILATION

GPU/CUDA: Compile

- Set the environment in the same way you set it for PGI or Intel
- Then compile the code
- Example below is for matrix multiplication code on another HPC machine

```
[$] module purge
[$] which nvcc
/usr/bin/which: no nvcc in (/usr/lib64/qt-
3.3/bin:/usr/local/bin:/bin:/usr/bin:/usr/local/sbin:/usr/sbin:/sbin:/opt/sdsc/bin:/o
pt/sdsc/sbin:/opt/ibutils/bin:/usr/java/latest/bin:/opt/pdsh/bin:/opt/rocks/bin:/opt/
rocks/sbin:/home/user/bin)
[$] module load cuda
[$] which nvcc
/usr/local/cuda-7.0/bin/nvcc
[$] nvcc -o matmul -I. matrixMul.cu
```

GPU/CUDA: check node for GPU card

Note: you will be able to [compile GPU](#) code on the login nodes, but they will not run.
To see if your node has GPU hardware, run `lspci`. Typically, login nodes do not have GPU hardware.

```
[$] lspci | grep VGA  
09:00.0 VGA compatible controller: ASPEED Technology, Inc. ASPEED Graphics Family  
(rev 30)
```

If the node **does** have a GPU card, you will see output similar to the following:

```
[$] ssh gpu-hotel "/sbin/lspci | grep VGA"  
01:00.0 VGA compatible controller: NVIDIA Corp.. NV44 [GeForce 6200 LE] (rev a1)  
02:00.0 VGA compatible controller: NVIDIA Corp.. GF100 [GeForce GTX 480] (rev a3)  
03:00.0 VGA compatible controller: NVIDIA Corp.. GF100 [GeForce GTX 480] (rev a3)
```

GPU/CUDA: Check Environment!

obtaining device information: enum_gpu.cu (1)

```
#include "../common/book.h"
int main( void ) {
    cudaDeviceProp prop;
    int count;
    HANDLE_ERROR( cudaGetDeviceCount( &count ) );
    for (int i=0; i< count; i++) {
        HANDLE_ERROR( cudaGetDeviceProperties( &prop, i ) );
        printf( " --- General Information for device %d ---\n", i );
        printf( "Name: %s\n", prop.name );
        printf( "Compute capability: %d.%d\n", prop.major, prop.minor );
        printf( "Clock rate: %d\n", prop.clockRate );
        printf( "Device copy overlap: " );
        if (prop.deviceOverlap)
            printf( "Enabled\n" );
        else
            printf( "Disabled\n" );
        printf( "Kernel execution timeout : " );
        if (prop.kernelExecTimeoutEnabled)
            printf( "Enabled\n" );
        else
            printf( "Disabled\n" );
    }
}
```

GPU/CUDA: Check Environment!

obtaining device information: enum_gpu.cu (2)

```
printf( "    --- Memory Information for device %d ---\n", i );
printf( "Total global mem:  %ld\n", prop.totalGlobalMem );
printf( "Total constant Mem:  %ld\n", prop.totalConstMem );
printf( "Max mem pitch:  %ld\n", prop.memPitch );
printf( "Texture Alignment:  %ld\n", prop.textureAlignment );
printf( "    --- MP Information for device %d ---\n", i );
printf( "Multiprocessor count:  %d\n",
        prop.multiProcessorCount );
printf( "Shared mem per mp:  %ld\n", prop.sharedMemPerBlock );
printf( "Registers per mp:  %d\n", prop.regsPerBlock );
printf( "Threads in warp:  %d\n", prop.warpSize );

printf( "Max threads per block:  %d\n",
        prop.maxThreadsPerBlock );
printf( "Max thread dimensions:  (%d, %d, %d)\n",
        prop.maxThreadsDim[0], prop.maxThreadsDim[1],
        prop.maxThreadsDim[2] );
printf( "Max grid dimensions:  (%d, %d, %d)\n",
        prop.maxGridSize[0], prop.maxGridSize[1],
        prop.maxGridSize[2] );
printf( "\n" );
}
```

GPU/CUDA: Check Environment*

```
-----  
--- General Information for device 0 ---  
Name: Tesla C1060  
Compute capability: 1.3  
Clock rate: 1296000  
Device copy overlap: Enabled  
Kernel execution timeout : Disabled  
--- Memory Information for device 0 ---  
Total global mem: 4294770688  
Total constant Mem: 65536  
Max mem pitch: 2147483647  
Texture Alignment: 256  
--- MP Information for device 0 ---  
Multiprocessor count: 30  
Shared mem per mp: 16384  
Registers per mp: 16384  
Threads in warp: 32  
Max threads per block: 512  
Max thread dimensions: (512, 512, 64)  
Max grid dimensions: (65535, 65535, 1)  
-----  
--- General Information for device 2 ---  
Name: GeForce GT 240  
Compute capability: 1.2  
Clock rate: 1340000  
Device copy overlap: Enabled  
Kernel execution timeout : Disabled  
--- Memory Information for device 2 ---  
Total global mem: 1073020928  
Total constant Mem: 65536  
Max mem pitch: 2147483647  
Texture Alignment: 256
```

```
-----  
--- General Information for device 1 ---  
Name: Tesla C1060  
Compute capability: 1.3  
Clock rate: 1296000  
Device copy overlap: Enabled  
Kernel execution timeout : Disabled  
--- Memory Information for device 1 ---  
Total global mem: 4294770688  
Total constant Mem: 65536  
Max mem pitch: 2147483647  
Texture Alignment: 256  
--- MP Information for device 1 ---  
Multiprocessor count: 30  
Shared mem per mp: 16384  
Registers per mp: 16384  
Threads in warp: 32  
Max threads per block: 512  
Max thread dimensions: (512, 512, 64)  
Max grid dimensions: (65535, 65535, 1)  
-----  
--- MP Information for device 2 ---  
Multiprocessor count: 12  
Shared mem per mp: 16384  
Registers per mp: 16384  
Threads in warp: 32  
Max threads per block: 512  
Max thread dimensions: (512, 512, 64)  
Max grid dimensions: (65535, 65535, 1)
```

*Data above is typical of GPUs, but not necessarily those on TSCC resources